

Ph.D. Thesis Defense

Generalized t-SNE and Beyond: Probabilistic Methods for Dimensionality
Reduction, Combinatorial Optimization, and Machine Learning

Yi Gu

Advisor: Antonio Auffinger

Department of Mathematics, Northwestern University

June 6, 2025

This thesis defense is about probabilistic algorithms along with the problems they aim to tackle.

- Two basic problems: Dimensionality and Dependence
- Dimensionality Reduction and t-SNE
- Dependence, Probabilistic Graphic Models and Lovász Local Lemma
- Autoencoder and Variational Factor Analysis

Dimensionality has been with us since the dawn of time. It refers to different aspects of the information. The scale of dimensionality drastically increases during the modern era because of advancements in mathematics and computer science.

- Increasing dimensionality means we can store more data, even the raw information of the process we aim to analyze.
- However, it does not mean we are automatically able to extract insights from the data.


The curse of dimensionality: a series of phenomena that arises in high-dimensional spaces that are usually not present or prominent in low-dimensional settings.

Example: The volume of a hypercube gets concentrated in its corners. Let C_d , S_d be the volume of the hypercube and hypersphere respectively.

$$\lim_{d \rightarrow \infty} \frac{S_d}{C_d} = \lim_{d \rightarrow \infty} \frac{\pi^{d/2}}{d 2^{d-1} \Gamma(d/2)} = 0. \quad (1)$$

Curse of dimensionality (cont.)

Example: the distribution of distance between uniformly distributed points is concentrated in high-dimensional spaces.



t-distributed Stochastic Neighbor Embedding (t-SNE) is an algorithm that find a low-dimensional representation of high-dimensional data.

$d, s \geq 2$ with $d > s$, and $n \geq 1$.

- **Input:** $X_1, \dots, X_n \in \mathbb{R}^d$ the original data points.
- **Output:** $Y_1, \dots, Y_n \in \mathbb{R}^s$ the low-dimensional representation.

Overall idea: KL-divergence (relative entropy) can be used to measure the difference between two distributions. t-SNE construct two distributions from both input and initial output data points, then minimize the KL-divergence by adjusting the output.

The concept of kernels

Definition (Generalized Input Kernel)

The input kernel $\mathbf{Ker_I}(x, y, \sigma) : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}^+$ is a function of the input data points and parameter σ such that for any $x, y \in \mathbb{R}^d$ and $\sigma > 0$,

$$\mathbf{Ker_I}(x, y, \sigma) = \mathbf{Ker_I}(y, x, \sigma). \quad (2)$$

Definition (Generalized Output Kernel)

The output kernel $\mathbf{Ker_O}(z, w) : \mathbb{R}^s \times \mathbb{R}^s \rightarrow \mathbb{R}^+$ is a function of the output data points such that for any $z, w \in \mathbb{R}^s$,

$$\mathbf{Ker_O}(z, w) = \mathbf{Ker_O}(w, z). \quad (3)$$

The concept of kernels(cont.)

Construct probability distributions from the input kernels:

$$p_{j|i} = \frac{\mathbf{Ker_I}(X_i, X_j, \sigma_i)}{\sum_{k \neq i} \mathbf{Ker_I}(X_i, X_k, \sigma_i)}, \quad p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}. \quad (4)$$

For output kernels:

$$q_{ij} = \frac{\mathbf{Ker_O}(Y_i, Y_j)}{\sum_{k \neq l} \mathbf{Ker_O}(Y_k, Y_l)}. \quad (5)$$

Original t-SNE can be recovered by letting

$$\mathbf{Ker_I}(x, y, \sigma) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right), \quad \mathbf{Ker_O}(z, w) = \frac{1}{1 + \|z - w\|^2}. \quad (6)$$

Recall

$$p_{j|i} = \frac{\mathbf{Ker_I}(X_i, X_j, \sigma_i)}{\sum_{k \neq i} \mathbf{Ker_I}(X_i, X_k, \sigma_i)}. \quad (7)$$

σ_i satisfies that for each $i = 1, \dots, n$,

$$- \sum_{j=1, j \neq i}^n p_{j|i} \log p_{j|i} = \log \mathbf{Perp} =: \log(n\rho), \rho \in (0, 1). \quad (8)$$

Intuition: σ_i is the zoom slide bar for the input data points.

- **Input:** $X_1, \dots, X_n \in \mathbb{R}^d \implies p_{ij}$.
- **Output:** $Y_1, \dots, Y_n \in \mathbb{R}^s \implies q_{ij}$.
- **Variables Involved:** p_{ij} are fixed. q_{ij} can be changed and is a function of Y_i and Y_j .
- **Loss Function (Goal):** minimize the KL-divergence:

$$L_{n,\rho}(X, Y) = \sum_{1 \leq i \neq j \leq n} p_{ij} \log \left(\frac{p_{ij}}{q_{ij}} \right). \quad (9)$$

- **End Result:**

$$Y^* = \arg \min_{Y \in (\mathbb{R}^s)^n} L_n(X, Y). \quad (10)$$

Continuous version of the entropy:

$$F_{\rho,\mu}(x, \sigma) = \int \frac{\mathbf{Ker_I}(x, x', \sigma)}{\int \mathbf{Ker_I}(x, x', \sigma) d\mu(x')} \log \left(\frac{\mathbf{Ker_I}(x, x', \sigma)}{\int \mathbf{Ker_I}(x, x', \sigma) d\mu(x')} \right) d\mu(x') + \log \rho. \quad (11)$$

$\sigma_{\rho,\mu}^* : \mathbb{R}^d \rightarrow \mathbb{R}$ the unique solution of the equation

$$F_{\rho,\mu}(x, \sigma) = 0. \quad (12)$$

σ_i serves the anchor point of distributions induced by the input data. $\sigma_{\rho,\mu}^*$ is its asymptotic limit.

Continuous version of induced probability distributions:

$\psi : \mathbb{R}^d \rightarrow \mathbb{R}$ be a given function and $p_\psi : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$ be defined by

$$p_\psi(x, x') = \frac{1}{2} \left(\frac{\mathbf{Ker_I}(x, x', \psi(x))}{\int \mathbf{Ker_I}(x, x', \psi(x)) d\mu(x')} + \frac{\mathbf{Ker_I}(x, x', \psi(x'))}{\int \mathbf{Ker_I}(x, x', \psi(x')) d\mu(x)} \right). \quad (13)$$

$q : \mathbb{R}^s \times \mathbb{R}^s \rightarrow \mathbb{R}$ be given by

$$q(y, y') = \frac{\mathbf{Ker_O}(y, y')}{\iint \mathbf{Ker_O}(y, y') d\mu(y) d\mu(y')}. \quad (14)$$

Continuous version of the loss function:

$$I_\rho(\mu) = \iint p_{\sigma_{\rho,\mu}^*}(x, x') \log \left(\frac{p_{\sigma_{\rho,\mu}^*}(x, x')}{q(y, y')} \right) d\mu(x, y) d\mu(x', y'). \quad (15)$$

Recall $\mu_X \in \mathcal{M}(\mathbb{R}^d)$ is an arbitrary given probability measure. We set $\mathcal{P}_X$ to be the set of probability measures on $\mathbb{R}^d \times \mathbb{R}^s$ whose marginal coincides with μ_X . Formally speaking,

$$\mathcal{P}_X := \{\mu \in \mathcal{M}(\mathbb{R}^d \times \mathbb{R}^s) : \text{the } \mathbb{R}^d \text{ marginal of } \mu \text{ is } \mu_X\}. \quad (16)$$

Lastly, define

$$\tilde{\mathcal{P}}_X := \left\{ \mu \in \mathcal{P}_X : \int (p_{\sigma_{\rho,\mu}^*}(x, x') - q(y, y')) \frac{\frac{\partial}{\partial y} \mathbf{Ker}_O(y, y')}{\mathbf{Ker}_O(y, y')} d\mu(x', y') = 0, \mu(x, y)\text{-a.e.} \right\}. \quad (17)$$

Condition (Input Kernel Conditions)

Let $w : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}_{\geq 0}$ be a fixed function. The input kernel has the form of

$$\mathbf{Ker_I}(x, x', \sigma) = \exp(-w(\sigma \|x - x'\|^\theta)) \quad (18)$$

Here, $\theta > 0$ is a fixed constant. w satisfies:

1. $\forall t \geq 0, w'(t) > 0. \lim_{t \rightarrow \infty} w(t) = \infty.$
2. $w'(t) + tw''(t) \geq 0$ for all $t \geq 0.$
3. The following integral converges:

$$\int_0^\infty t^{d-1} w(t^\theta) \exp(-w(t^\theta)) dt < \infty. \quad (19)$$

Since the all probability masses are normalized, we can assume without loss of generality that $w(0) = 0$ from now on.

Condition (Output Kernel Conditions)

The output kernel depends only on the distance between two output data points y and y' . In particular, $\mathbf{Ker}_O(y, y') = k(\|y - y'\|)$ and:

- 1. $k : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}^+$ is decreasing and bounded.*
- 2. $\mathbf{Ker}_O$ satisfies*

$$\iint_{(0, \infty)^2} \mathbf{Ker}_O(y, y') dy dy' < \infty. \quad (20)$$

- 3. k has bounded derivative.*
- 4. $k'(0) = 0$.*

Theorem

Assume μ_X has a continuous density $f(x)$ and compact support. Let $X_i \in \mathbb{R}^d$ be i.i.d random variables with probability measure μ_X . If condition 1 and condition 2 hold, then for any $\rho \in (0, 1)$,

$$\lim_{n \rightarrow \infty} \inf_Y L_{n,\rho}(X, Y) = \inf_{\mu \in \tilde{\mathcal{P}}_X} I_\rho(\mu). \quad (21)$$

Moreover, there exist a measure μ^ and a sub-sequence $\{n_k\}$ such that $\mu_{n_k} \rightarrow \mu^*$ weakly and $I_\rho(\mu^*) = \inf_{\mu \in \tilde{\mathcal{P}}_X} I_\rho(\mu)$.*

Theorem

Under the same assumptions of theorem 3, μ^ has compact support.*

The problems of dependence

